

Assignment1.1 Sutow Brett

June 8, 2021

```
[1]: #Example One#
from tensorflow import keras
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.optimizers import RMSprop
```

```
[2]: batch_size = 128
num_classes = 10
epochs = 20
```

```
[3]: # the data, split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

x_train = x_train.reshape(60000, 784)
x_test = x_test.reshape(10000, 784)
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')
```

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-
datasets/mnist.npz
11493376/11490434 [=====] - 1s 0us/step
60000 train samples
10000 test samples
```

```
[4]: # convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

model = Sequential()
model.add(Dense(512, activation='relu', input_shape=(784,)))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.2))
```

```

model.add(Dense(num_classes, activation='softmax'))

model.summary()

model.compile(loss='categorical_crossentropy',
              optimizer=RMSprop(),
              metrics=['accuracy'])

history = model.fit(x_train, y_train,
                   batch_size=batch_size,
                   epochs=epochs,
                   verbose=1,
                   validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 512)	401920
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 512)	262656
dropout_1 (Dropout)	(None, 512)	0
dense_2 (Dense)	(None, 10)	5130

Total params: 669,706

Trainable params: 669,706

Non-trainable params: 0

Epoch 1/20

469/469 [=====] - 5s 11ms/step - loss: 0.4359 - accuracy: 0.8619 - val_loss: 0.1130 - val_accuracy: 0.9647

Epoch 2/20

469/469 [=====] - 4s 9ms/step - loss: 0.1104 - accuracy: 0.9670 - val_loss: 0.0888 - val_accuracy: 0.9726

Epoch 3/20

469/469 [=====] - 4s 9ms/step - loss: 0.0727 - accuracy: 0.9775 - val_loss: 0.0709 - val_accuracy: 0.9792

Epoch 4/20

469/469 [=====] - 4s 9ms/step - loss: 0.0574 - accuracy: 0.9826 - val_loss: 0.0847 - val_accuracy: 0.9772

Epoch 5/20
469/469 [=====] - 4s 9ms/step - loss: 0.0464 -
accuracy: 0.9855 - val_loss: 0.0785 - val_accuracy: 0.9814

Epoch 6/20
469/469 [=====] - 4s 9ms/step - loss: 0.0409 -
accuracy: 0.9880 - val_loss: 0.0795 - val_accuracy: 0.9796

Epoch 7/20
469/469 [=====] - 4s 9ms/step - loss: 0.0364 -
accuracy: 0.9899 - val_loss: 0.0822 - val_accuracy: 0.9816

Epoch 8/20
469/469 [=====] - 4s 9ms/step - loss: 0.0327 -
accuracy: 0.9905 - val_loss: 0.0865 - val_accuracy: 0.9830

Epoch 9/20
469/469 [=====] - 4s 9ms/step - loss: 0.0281 -
accuracy: 0.9913 - val_loss: 0.0799 - val_accuracy: 0.9826

Epoch 10/20
469/469 [=====] - 4s 9ms/step - loss: 0.0248 -
accuracy: 0.9929 - val_loss: 0.0844 - val_accuracy: 0.9836

Epoch 11/20
469/469 [=====] - 4s 9ms/step - loss: 0.0237 -
accuracy: 0.9930 - val_loss: 0.0879 - val_accuracy: 0.9842

Epoch 12/20
469/469 [=====] - 4s 9ms/step - loss: 0.0232 -
accuracy: 0.9932 - val_loss: 0.0978 - val_accuracy: 0.9802

Epoch 13/20
469/469 [=====] - 4s 9ms/step - loss: 0.0238 -
accuracy: 0.9934 - val_loss: 0.0949 - val_accuracy: 0.9833

Epoch 14/20
469/469 [=====] - 4s 9ms/step - loss: 0.0191 -
accuracy: 0.9948 - val_loss: 0.0899 - val_accuracy: 0.9829

Epoch 15/20
469/469 [=====] - 4s 9ms/step - loss: 0.0185 -
accuracy: 0.9946 - val_loss: 0.1074 - val_accuracy: 0.9821

Epoch 16/20
469/469 [=====] - 4s 9ms/step - loss: 0.0179 -
accuracy: 0.9950 - val_loss: 0.1126 - val_accuracy: 0.9824

Epoch 17/20
469/469 [=====] - 4s 9ms/step - loss: 0.0168 -
accuracy: 0.9956 - val_loss: 0.1350 - val_accuracy: 0.9826

Epoch 18/20
469/469 [=====] - 4s 9ms/step - loss: 0.0186 -
accuracy: 0.9950 - val_loss: 0.1266 - val_accuracy: 0.9823

Epoch 19/20
469/469 [=====] - 4s 9ms/step - loss: 0.0173 -
accuracy: 0.9953 - val_loss: 0.1194 - val_accuracy: 0.9822

Epoch 20/20
469/469 [=====] - 4s 9ms/step - loss: 0.0144 -
accuracy: 0.9957 - val_loss: 0.1305 - val_accuracy: 0.9832

Test loss: 0.1305491179227829
Test accuracy: 0.9832000136375427

```
[6]: #Example Two#  
import sys  
from random import random  
from operator import add  
  
from pyspark.sql import SparkSession
```

```
[14]: {  
  "cells": [  
    {  
      "cell_type": "code",  
      "execution_count": 1,  
      "id": "communist-label",  
      "metadata": {},  
      "outputs": [],  
      "source": [  
        "#Example One#\n",  
        "from tensorflow import keras\n",  
        "from tensorflow.keras.datasets import mnist\n",  
        "from tensorflow.keras.models import Sequential\n",  
        "from tensorflow.keras.layers import Dense, Dropout\n",  
        "from tensorflow.keras.optimizers import RMSprop\n"  
      ]  
    },  
    {  
      "cell_type": "code",  
      "execution_count": 2,  
      "id": "aboriginal-reproduction",  
      "metadata": {},  
      "outputs": [],  
      "source": [  
        "batch_size = 128\n",  
        "num_classes = 10\n",  
        "epochs = 20\n"  
      ]  
    },  
    {  
      "cell_type": "code",  
      "execution_count": 3,  
      "id": "mediterranean-wagner",  
      "metadata": {},  
      "outputs": [  
        {  
          "name": "stdout",  

```

```

    "output_type": "stream",
    "text": [
        "Downloading data from https://storage.googleapis.com/tensorflow/
↳tf-keras-datasets/mnist.npz\n",
        "11493376/11490434 [=====] - 1s 0us/step\n",
        "60000 train samples\n",
        "10000 test samples\n"
    ]
}
],
"source": [
    "# the data, split between train and test sets\n",
    "(x_train, y_train), (x_test, y_test) = mnist.load_data()\n",
    "\n",
    "x_train = x_train.reshape(60000, 784)\n",
    "x_test = x_test.reshape(10000, 784)\n",
    "x_train = x_train.astype('float32')\n",
    "x_test = x_test.astype('float32')\n",
    "x_train /= 255\n",
    "x_test /= 255\n",
    "print(x_train.shape[0], 'train samples')\n",
    "print(x_test.shape[0], 'test samples')\n"
]
},
{
    "cell_type": "code",
    "execution_count": 4,
    "id": "second-jason",
    "metadata": {},
    "outputs": [
        {
            "name": "stdout",
            "output_type": "stream",
            "text": [
                "Model: \"sequential\"\n",
                "-----\n",


| Layer (type)        | Output Shape | Param # |
|---------------------|--------------|---------|
| dense (Dense)       | (None, 512)  | 401920  |
| dropout (Dropout)   | (None, 512)  | 0       |
| dense_1 (Dense)     | (None, 512)  | 262656  |
| dropout_1 (Dropout) | (None, 512)  | 0       |
| dense_2 (Dense)     | (None, 10)   | 5130    |


```

```

"=====\n",
"Total params: 669,706\n",
"Trainable params: 669,706\n",
"Non-trainable params: 0\n",
"-----\n",
"Epoch 1/20\n",
"469/469 [=====] - 5s 11ms/step - loss: 0.4359 -  

↳accuracy: 0.8619 - val_loss: 0.1130 - val_accuracy: 0.9647\n",
"Epoch 2/20\n",
"469/469 [=====] - 4s 9ms/step - loss: 0.1104 -  

↳accuracy: 0.9670 - val_loss: 0.0888 - val_accuracy: 0.9726\n",
"Epoch 3/20\n",
"469/469 [=====] - 4s 9ms/step - loss: 0.0727 -  

↳accuracy: 0.9775 - val_loss: 0.0709 - val_accuracy: 0.9792\n",
"Epoch 4/20\n",
"469/469 [=====] - 4s 9ms/step - loss: 0.0574 -  

↳accuracy: 0.9826 - val_loss: 0.0847 - val_accuracy: 0.9772\n",
"Epoch 5/20\n",
"469/469 [=====] - 4s 9ms/step - loss: 0.0464 -  

↳accuracy: 0.9855 - val_loss: 0.0785 - val_accuracy: 0.9814\n",
"Epoch 6/20\n",
"469/469 [=====] - 4s 9ms/step - loss: 0.0409 -  

↳accuracy: 0.9880 - val_loss: 0.0795 - val_accuracy: 0.9796\n",
"Epoch 7/20\n",
"469/469 [=====] - 4s 9ms/step - loss: 0.0364 -  

↳accuracy: 0.9899 - val_loss: 0.0822 - val_accuracy: 0.9816\n",
"Epoch 8/20\n",
"469/469 [=====] - 4s 9ms/step - loss: 0.0327 -  

↳accuracy: 0.9905 - val_loss: 0.0865 - val_accuracy: 0.9830\n",
"Epoch 9/20\n",
"469/469 [=====] - 4s 9ms/step - loss: 0.0281 -  

↳accuracy: 0.9913 - val_loss: 0.0799 - val_accuracy: 0.9826\n",
"Epoch 10/20\n",
"469/469 [=====] - 4s 9ms/step - loss: 0.0248 -  

↳accuracy: 0.9929 - val_loss: 0.0844 - val_accuracy: 0.9836\n",
"Epoch 11/20\n",
"469/469 [=====] - 4s 9ms/step - loss: 0.0237 -  

↳accuracy: 0.9930 - val_loss: 0.0879 - val_accuracy: 0.9842\n",
"Epoch 12/20\n",
"469/469 [=====] - 4s 9ms/step - loss: 0.0232 -  

↳accuracy: 0.9932 - val_loss: 0.0978 - val_accuracy: 0.9802\n",
"Epoch 13/20\n",
"469/469 [=====] - 4s 9ms/step - loss: 0.0238 -  

↳accuracy: 0.9934 - val_loss: 0.0949 - val_accuracy: 0.9833\n",
"Epoch 14/20\n",

```

```

    "469/469 [=====] - 4s 9ms/step - loss: 0.0191 -_
    ↳accuracy: 0.9948 - val_loss: 0.0899 - val_accuracy: 0.9829\n",
    "Epoch 15/20\n",
    "469/469 [=====] - 4s 9ms/step - loss: 0.0185 -_
    ↳accuracy: 0.9946 - val_loss: 0.1074 - val_accuracy: 0.9821\n",
    "Epoch 16/20\n",
    "469/469 [=====] - 4s 9ms/step - loss: 0.0179 -_
    ↳accuracy: 0.9950 - val_loss: 0.1126 - val_accuracy: 0.9824\n",
    "Epoch 17/20\n",
    "469/469 [=====] - 4s 9ms/step - loss: 0.0168 -_
    ↳accuracy: 0.9956 - val_loss: 0.1350 - val_accuracy: 0.9826\n",
    "Epoch 18/20\n",
    "469/469 [=====] - 4s 9ms/step - loss: 0.0186 -_
    ↳accuracy: 0.9950 - val_loss: 0.1266 - val_accuracy: 0.9823\n",
    "Epoch 19/20\n",
    "469/469 [=====] - 4s 9ms/step - loss: 0.0173 -_
    ↳accuracy: 0.9953 - val_loss: 0.1194 - val_accuracy: 0.9822\n",
    "Epoch 20/20\n",
    "469/469 [=====] - 4s 9ms/step - loss: 0.0144 -_
    ↳accuracy: 0.9957 - val_loss: 0.1305 - val_accuracy: 0.9832\n",
    "Test loss: 0.1305491179227829\n",
    "Test accuracy: 0.9832000136375427\n"
]
}
],
"source": [
    "# convert class vectors to binary class matrices\n",
    "y_train = keras.utils.to_categorical(y_train, num_classes)\n",
    "y_test = keras.utils.to_categorical(y_test, num_classes)\n",
    "\n",
    "model = Sequential()\n",
    "model.add(Dense(512, activation='relu', input_shape=(784,)))\n",
    "model.add(Dropout(0.2))\n",
    "model.add(Dense(512, activation='relu'))\n",
    "model.add(Dropout(0.2))\n",
    "model.add(Dense(num_classes, activation='softmax'))\n",
    "\n",
    "model.summary()\n",
    "\n",
    "model.compile(loss='categorical_crossentropy',\n",
    "               optimizer=RMSprop(),\n",
    "               metrics=['accuracy'])\n",
    "\n",
    "history = model.fit(x_train, y_train,\n",
    "                    batch_size=batch_size,\n",
    "                    epochs=epochs,

```

```

        verbose=1,\n",
        validation_data=(x_test, y_test))\n",
"score = model.evaluate(x_test, y_test, verbose=0)\n",
"print('Test loss:', score[0])\n",
"print('Test accuracy:', score[1])"
]
},
{
"cell_type": "code",
"execution_count": 6,
"id": "confidential-playback",
"metadata": {},
"outputs": [],
"source": [
"#Example Two#\n",
"import sys\n",
"from random import random\n",
"from operator import add\n",
"\n",
"from pyspark.sql import SparkSession\n"
]
},
{
"cell_type": "code",
"execution_count": 14,
"id": "academic-leeds",
"metadata": {},
"outputs": [
{
"name": "stdout",
"output_type": "stream",
"text": [
"Pi is roughly 3.142760\n"
]
}
]
},
"source": [
"\n",
"if __name__ == \"__main__\":\n",
"    \"\"\"\n",
"        Usage: pi [partitions]\n",
"    \"\"\"\n",
"    spark = SparkSession\\\n",
"        .builder\\\n",
"        .appName(\"PythonPi\")\\\n",
"        .getOrCreate()\n",
"\n"
]

```



```

    "    partitions = 2 #int(sys.argv[1]) if len(sys.argv) > 1 else 2#\n",
    "    n = 100000 * partitions\n",
    "\n",
    "    def f(_):\n",
    "        x = random() * 2 - 1\n",
    "        y = random() * 2 - 1\n",
    "        return 1 if x ** 2 + y ** 2 <= 1 else 0\n",
    "\n",
    "    count = spark.sparkContext.parallelize(range(1, n + 1), partitions).\n
↪map(f).reduce(add)\n",
    "    print(\"Pi is roughly %f\" % (4.0 * count / n))\n",
    "\n",
    "    spark.stop()"
]
},
{
    "cell_type": "code",
    "execution_count": null,
    "id": "desperate-resident",
    "metadata": {},
    "outputs": [],
    "source": []
}
],
"metadata": {
    "kernel_spec": {
        "display_name": "Python 3",
        "language": "python",
        "name": "python3"
    },
    "language_info": {
        "codemirror_mode": {
            "name": "ipython",
            "version": 3
        },
        "file_extension": ".py",
        "mimetype": "text/x-python",
        "name": "python",
        "nbconvert_exporter": "python",
        "pygments_lexer": "ipython3",
        "version": "3.8.8"
    }
},
"nbformat": 4,
"nbformat_minor": 5
}

```

```

if __name__ == "__main__":
    """
        Usage: pi [partitions]
    """
    spark = SparkSession\
        .builder\
        .appName("PythonPi")\
        .getOrCreate()

    partitions = 2 #int(sys.argv[1]) if len(sys.argv) > 1 else 2#
    n = 100000 * partitions

    def f(_):
        x = random() * 2 - 1
        y = random() * 2 - 1
        return 1 if x ** 2 + y ** 2 <= 1 else 0

    count = spark.sparkContext.parallelize(range(1, n + 1), partitions).map(f).\
    ↪reduce(add)
    print("Pi is roughly %f" % (4.0 * count / n))

    spark.stop()

```

Pi is roughly 3.142760

[]: